




# Graph Databases Use Cases






# What's a Graph?



# Graph data model



# Relational Tables


## SQL Join Hell (1)

Customer		
<b>Id</b>	Name	<b>Address</b>
1	Robert	3
2	Lars	7
3	Michael	23




1:1 Relationship

Customer		Address	
<b>Id</b>	Name	<b>Customer</b>	Location
1	Robert	3	Berlin
2	Lars	4	Munich
3	Michael	7	Dresden
		23	Leipzig




1:n Relationship

Customer		CustomerAddress		Address	
<b>Id</b>	Name	<b>CId</b>	<b>AId</b>	<b>Id</b>	Location
1	Robert	1	3	3	Berlin
2	Lars	2	7	7	Dresden
3	Michael	2	8	8	New York
		3	23	23	Leipzig




m:n Relationship

# Join this way...




```
SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
```





```
SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
WHERE B.key IS NULL
```


```
SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
```



```
SELECT <auswahl>
FROM tabelleA A
INNER JOIN tabelleB B
ON A.key = B.key
```




```
SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
```




```
SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
```

```
SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
OR B.key IS NULL
```



# The Problem


- all JOINs are executed **every time you query (traverse) the relationship**
- executing a JOIN means to **search for a key in another table**
- with Indices executing a JOIN means to **lookup a key**
- B-Tree Index:  $O(\log(n))$
- more entries => more lookups => slower JOINs




**People**

**Attend**


**Conferences**



# A Property Graph



# The Neo4j Secret Sauce



- Pointers instead of look-ups
- Do all your “Joining” on creation
- Spin spin spin through this data structure


# Graph Buzz!

MacArthur 'Genius Grant' Winner Maria Chudnovsky on Graph Theory

Wednesday, October 03, 2012




## Bright Launches Bright Packed With Jobs Data Seeking Tips



## Facebook's Social Graph, Neo4j show rising use of graph databases

**Summary:** Facebook's Social Graph -- the database underlying its Graph Search engine unveiled yesterday-- is just one of many graph databases being employed for complex, connected data. Neo4j



259

Like

969





Tweet

48

+1

173

Share







I saw my own Interest Graph and it's scary-accurate. We'd certainly pay for the ability to use the Gravity personalization technology I saw today at TechCrunch to help target content to users.

TC Michael Arrington, TechCrunch

# The Neo4j Graph Database

- Neo4j is the **leading graph database** in the world today
  - Most widely deployed: **500,000+** downloads
  - Largest ecosystem: active forums, code contributions, etc
  - Most mature product: in development since 2000, in **24/7 production** since **2003**



neo4j leaves me speechless. Good job at building the best graph database in the world! ★

2:41 AM Mar 24th via Tweetie  
Retweeted by you and 1 other

Reply Retweeted (Undo)

 **fokussiertnet**  
Daniel

If #Cassandra rules its league (#distributed #decentralized #ColumnFamily) then #Neo4j does it in its own, which is #GraphDB. #noSQL ★

3:12 AM Apr 8th via TweetDeck  
Retweeted by you



Reply Retweeted (Undo)

 **alisohani**  
Ali Sohani

# Early Adopters of Graph Tech



facebook.



Google™



# Survival of the Fittest

## Evolution of Web Search

### Pre-1999

#### WWW Indexing

**Index**

Page numbers in bold then refer to key term definitions  
 Page numbers in italics refer to images or diagrams  
 Page numbers followed by a “t” indicate a table

**A**


- absolute temperature scale, 380–381
- absolute zero, 381
- acceleration of gravity, A.23t
- accuracy, A.5
- acetic acid ( $\text{CH}_3\text{COOH}$ )

  - buffer, 575–576, 581–582
  - conjugate acid-base pairs, 540
  - ionization constant, 553, 558
  - manufacture of, 451
  - titration, 590–592
  - as weak acid, 148, 149, 510–512
  - acid-base pairs, conjugate, 848–849
  - acid-base reactions, 538
  - ionization of water, 540–542
  - gas-forming exchange, 550–551
  - net ionic equations for, 148–150
  - nernst equation, 146–150, 261–266
  - of salts, 148–151, 561–566
  - acid-base titrations, 597–598
  - amide solutions, 846, 850–852, A.32n–A.33t
  - and ionization constant, 850–851
  - and ionization constant ( $K_{\text{a}}$ ). *See* ionization constants, acids ( $K_{\text{a}}$ )
  - anhydrides, 143. *See also* acid-base reactions, ionization constants, acidic ( $K_{\text{a}}$ ), specific esters, e.g., carboxylic acids, lacto acids
  - Brensted-Lowry concept, 533–544
  - buffer solutions, 575–586
  - conjugate acid-base pairs, 540–544
  - equilibrium constant, 47%
  - ionization constants, 510–511, A.23t
  - Lewis, 546–548
  - organic, 544
  - pH scale, 547–550
  - properties, 143–145
  - solubility of salts, 597–598
  - solutions, 846, 850–852, A.32n–A.33t
  - strengths, 143–146, 533–536
  - titrations, 597–598
  - water's role, 540
  - weak acids, 55, 210–211
  - activated complex, 433
  - activation energies ( $E_a$ ), 434, 438–440, 443, 447, 449–450
  - activation energy, 449
  - activation, 466, 247, 703–706
  - activity series, metals, 159–160
  - actual yields, 123
  - addition, 43, A.6, A.9
  - addition, significant figures in, A.6

Discrete Data 

### 1999 - 2012


#### Google Invents PageRank



Connected Data  
(Simple)

### 2012-?

#### Google Knowledge Graph, Facebook Graph Search



Connected Data  
(Rich)

# Open Source Example

Neo Graph Search    Home    Profile    Graph Search    Likes    Friends    Visualization    [Max De Marzi](#)

## Graph Search

friends who like Neo4j  Try: [friends who like wine](#) , [people who like wine and cheese](#) , [people who like cycling and live in Florida](#)

**Cypher Query:**

```
START me = node({me}), thing = node:things({thing})
MATCH me -[:friends]-> people, people -[:likes]-> thing
RETURN DISTINCT people , people.uid, people.name, people.image_url
LIMIT 100
Parameters: {"me":>1, "thing":>"name: Neo4j"}
```

Want your own Graph Search? Contact [me](#) and learn more about [Neo4j](#) and [NeoTechnology](#)



Peter Neubauer



Andres Taylor

<http://maxdemarzi.com/2013/01/28/facebook-graph-search-with-cypher-and-neo4j/>

# Survival of the Fittest

## Evolution of Online Recruiting

### 1999

#### Keyword Search

[Home](#) > **Job Search**

Search over **150,000** U.S. jobs.

Perform your search below or get [tips on searching](#).

#### Location Search:

----- Select a location -----

- Alaska-Anchorage
- Alaska-Fairbanks
- Alaska-Juneau
- Alaska-Valdez
- Alabama-Anniston

#### Category Search:

----- Select a category -----

- Accounting/Finance/Banking
- Administrative/Clerical
- Creative Arts/Media
- Education/Training
- Engineering/Architecture/Design

#### Keyword Search:

[Search Jobs](#) [Clear](#)

Discrete Data



### 2011-12

#### Social Discovery



#### Most jobs are found through an inside connection

Each friend that joins Glassdoor allows you to see more connections at more companies

9 friends on glassdoor	3,905 inside connections	3,882 companies
------------------------	--------------------------	-----------------

[Invite more friends](#) — ask them to share their connections

#### Jobs with Connections

[Sr. Statistical Analyst, Product Innovation](#)  
Netflix — Los Gatos, CA  
From: Job.com — 1 days ago

[Java Server Software Engineer](#)  
Electronic Arts — Redwood City, CA  
From: Experiteer — 3 days ago


Want better jobs? [Tell us your current job title](#)

[Creative Director](#)  
frog design — San Francisco, CA  
From: Experiteer — 8 days ago

[EMERGENCY MEDICAL TECH](#)  
U. S. NAVY — Fremont, CA  
From: Monster — 14 days ago

Connected Data

# Open Source Example



<http://maxdemarzi.com/2012/10/18/matches-are-the-new-hotness/>

# Open Source Example

```
START me=node:users_index(name={user})
MATCH skills<-[ :has]-me-[:lives_in]->city<-[ :in_location]-job-[:requires]->requirements
WHERE me-[:has]->()-<[:requires]-job
WITH DISTINCT city.name AS city_name,
      job.name AS job_name,
      LENGTH(me-[:has]->()-<[:requires]-job) AS matching_skills,
      LENGTH(job-[:requires]->()) AS job_requires,
      COLLECT(DISTINCT requirements.name) AS req_names,
      COLLECT(DISTINCT skills.name) AS skill_names
RETURN city_name, job_name,
       FILTER(name IN req_names WHERE NOT name IN skill_names) AS missing
ORDER BY matching_skills / job_requires DESC, job_requires
LIMIT 10
```

<http://maxdemarzi.com/2012/10/18/matches-are-the-new-hotness/>

# Open Source Example

## Matching Jobs for Daniel

Reload the page or click [Match Another User](#)

Job	Match	Missing
Spring-C-Django-SQL in <i>Dallas</i>	100% Match	
Neo4j-C-CSS in <i>Dallas</i>	Partial Match	<span>+ CSS</span>
C-Spring-Java in <i>Dallas</i>	Partial Match	<span>+ Java</span>
Java-Redis-C-Neo4j in <i>Dallas</i>	Partial Match	<span>+ Java</span> <span>+ Redis</span>


<http://maxdemarzi.com/2012/10/18/matches-are-the-new-hotness/>

# Emergent Graph in Other Industries




(Actual Neo4j Graphs)

## Content Management & Access Control




## Network Cell Analysis



## Bioinformatics




## Geo Routing (Public Transport)




## Insurance Risk Analysis



## Network Asset Management



# Open Source Example




<http://maxdemarzi.com/2013/03/18/permission-resolution-with-neo4j-part-1/>

# Emergent Graph in Other Industries


(Actual Neo4j Graphs)



## Web Browsing




## Portfolio Analytics



## Gene Sequencing



## Mobile Social Application



# Open Source Example

Neo Love Matching    Match Another User


## Matching Loves for Luvenia

Reload the page or click [Match Another User](#)

Love	Compatibility	Your Matches	Their Matches
Chet in Philadelphia	2.0	<input checked="" type="checkbox"/> Humorous <input checked="" type="checkbox"/> Thoughtful	<input checked="" type="checkbox"/> Likeable
Mckinley in Philadelphia	1.0	<input checked="" type="checkbox"/> Serious	<input checked="" type="checkbox"/> Creative
Gene in Philadelphia	1.0	<input checked="" type="checkbox"/> Serious	<input checked="" type="checkbox"/> Likeable
Keith in Philadelphia	1.0	<input checked="" type="checkbox"/> Quiet	<input checked="" type="checkbox"/> Inventive

<http://maxdemarzi.com/2013/04/19/match-making-with-neo4j/>

# Curriculum Graph



# Early Adopter Segments

(What we expected to happen - view from several years ago)



Core Industries & Use Cases:	Web / ISV	Finance & Insurance	Datacom / Telecom
Network & Data Center Management			
MDM			
Social			
Geo			


# Neo4j Adoption Snapshot

## Select Commercial Customers (Community Users Not Included)



Core Industries & Use Cases:	Software	Financial Services	Telecommunications	Web Social, HR & Recruiting	Health Care & Life Sciences	Media & Publishing	Energy, Services, Automotive, Gov't, Logistics, Education, Gaming, Other
Network & Data Center Management	Zenoss NetApp SERENA gen VIRTUAL INSTRUMENTS		hp SFR				Juice PLUS+ gonefinestay teachscape
MDM / System of Record			CISCO Deutsche Telekom maaii Let's connect		ZEPHYR HEALTH INC HealthUnlocked	LIFECHURCH.TV	
Social	Glowbl	ICE Global Markets in clear view			viadeo glassdoor SharePractice	SQUIDOO indiatimes	gameSYS Accenture Global 500 Logistics shuttle
Geo	DingLicom		Justdial India's No. 1 local search engine		classmates.com®		
Identity & Access Mgmt	aikux.com identropy	Global 500 Finance	telenor			zeebox LifeWay Biblical Solutions for Life	DOSB NEW MEDIA GMBH DEUTSCHER OLYMPISCHER SPORTBUND
Content Management	springcm Adobe			Dshini SHRM SOCIETY FOR HUMAN RESOURCE MANAGEMENT	SevenBridges genomics	dedbel <fuseworks/>	
Recommend-ations	LIQUID COMMON Social Web Mobile			hinge careerbuilder InfoJobs	Curaspan HEALTH GROUP	Perigee CHIP	research now compete
BI, CRM, Impact Analysis, Fraud Detection, Resource Optimization, etc.	AXON ACTIVE Focusing on your decisions kitedesk SODIFRANCE CONSEIL, TECHNOLOGIES & SERVICES IT idMISSION ENHANCING LIVES Humanvest.co     DRW TRADING GROUP	Global 500 Telcommunication		moviepilot t KLM van Kompagnie	janssen AmeriPharmaceutical Companies of Johnson & Johnson	DRAKER Impact Technologies A Sikorsky Innovations Company LOCKHEED MARTIN	Global 500 Energy Global 500 Aerospace

# What Can You Do With Graphs?






```
MATCH (me:Person) - [:IS_FRIEND_OF] -> (friend),  
(friend) - [:LIKES] -> (restaurant),  
(restaurant) - [:LOCATED_IN] -> (city:Location),  
(restaurant) - [:SERVES] -> (cuisine:Cuisine)
```


```
WHERE me.name = 'Philip' AND city.location='New York' AND  
cuisine.cuisine='Sushi'
```

```
RETURN restaurant.name
```



# Of course.. a graph is a graph is a graph

*What drugs will bind to protein X and not interact with drug Y?*



# Connected Query Performance

# Connected Query Performance



*Query Response Time\* =  
f(graph density, graph size, query degree)*

- **Graph density** (avg # rel's / node)
- **Graph size** (total # of nodes in the graph)
- **Query degree** (# of hops in one's query)

## RDBMS:

>> exponential slowdown as each factor increases


## Neo4j:

>> Performance remains constant as graph size increases

>> Performance slowdown is linear or better as density & degree increase

# Connected Query Performance

## RDBMS vs. Native Graph Database



# Graph db performance

- a sample social graph
  - with ~1,000 persons
- average 50 friends per person
- `pathExists(a,b)` limited to depth 4
- caches warmed up to eliminate disk I/O

Database	# persons	query time
MySQL		
Neo4j		
Neo4j		

# Graph db performance

- a sample social graph
  - with ~1,000 persons
- average 50 friends per person
- `pathExists(a,b)` limited to depth 4
- caches warmed up to eliminate disk I/O

Database	# persons	query time
MySQL	1,000	2,000 ms
Neo4j		
Neo4j		

# Graph db performance

- a sample social graph
  - with ~1,000 persons
- average 50 friends per person
- `pathExists(a,b)` limited to depth 4
- caches warmed up to eliminate disk I/O

Database	# persons	query time
MySQL	1,000	2,000 ms
Neo4j	1,000	2 ms
Neo4j		


# Graph db performance

- a sample social graph
  - with ~1,000 persons
- average 50 friends per person
- `pathExists(a,b)` limited to depth 4
- caches warmed up to eliminate disk I/O

Database	# persons	query time
MySQL	1,000	2,000 ms
Neo4j	1,000	2 ms
Neo4j	1,000,000	2 ms


\*Additional Third Party Benchmark Available in *Neo4j in Action*: <http://www.manning.com/partner/>

# The Zone of SQL Adequacy



# Graph Technology Ecosystem

# #1: Graph Local Queries



e.g. Recommendations, Friend-of-Friend, Shortest Path

# What is a Graph Database

“A **graph database**... is an online database management system with CRUD methods that expose a graph data model”<sup>1</sup>

- Two important properties:
  - **Native graph storage engine:** *written from the ground up to manage graph data*
  - **Native graph processing,** including *index-free adjacency* to facilitate traversals

# Graph Databases are Designed to:

1. Store inter-connected data
2. Make it easy to make sense of that data
3. Enable extreme-performance operations for:
  - Discovery of connected data patterns
  - Relatedness queries > depth 1
  - Relatedness queries of arbitrary length
4. Make it easy to evolve the database

# Top Reasons People Use Graph Databases

- 1. Problems with Join performance.**
- 2. Continuously evolving data set (often involves wide and sparse tables)**
- 3. The Shape of the Domain is naturally a graph**
- 4. Open-ended business requirements necessitating fast, iterative development.**


The Viadeo Group forms the world's second largest professional social network with over 50 million members.

- N°1 in France with over 7 million members,
- N°1 in China with over 14 millions members.
- Strong presence in Europe and in emerging markets, strong growth in South America
- Presence in Russia since December 2011 through a joint venture with leading media group : Sanoma Independent Media
- April 2012 : record funding of 24 million euros (FSI, historical shareholders and new entrants)
- Headquarters in Paris and local offices on 5 continents
- 400 employees




- The Viadeo Group consists of:

[ApnaCircle.com](http://ApnaCircle.com)[天際網  
www.tianji.com](http://www.tianji.com)



- In-house algorithm
- Network storage in **MySQL Database**

```
CREATE TABLE `Network` (
    `memberId` int(11) NOT NULL DEFAULT '0',
    `L1` mediumblob NOT NULL,
    `L2` mediumblob NOT NULL,
    PRIMARY KEY (`memberId`)
) ENGINE=InnoDB;
```



## LIMITATIONS

- 1) Important latency for complete update



- 2) Massive bandwidth impact for internal network

- 3) 48 hours to restart from scratch




## Update the network (old-fashioned style)

Member A and Member B are now in contacts


- Update of A.L1 + B.L1 and A.L2 + B.L2
- Retrieving A.L1 + B.L1 and update \*.L2

Example:

- A has 500 contacts
  - B has 150 contacts
- ➔  $500 + 150 + 2 = \textbf{652 updates!}$



## GRAPH DATABASE



Wait what?



New Users?  
Real Time Updates?

## BENEFITS



Very easy to integrate  
(less than 2 months)

Instantaneous  
graph updates




Backup /  
Restore



High Availability




# Graph Database Deployment




# Graph Dashboards

## The Power of Visualization

# Fraud Detection & Money Laundering




# IT Service Dependencies



# Working with Graphs

## Case Studies & Working Examples



# Cypher

## ASCII art Graph Patterns



```
MATCH (A) -[:LOVES]-> (B)  
WHERE A.name = "A"  
RETURN B as lover
```

# Social Example

# Practical Cypher


## Social Graph - Create



### CREATE

```
(joe:Person {name:"Joe"}) ,  
(bob:Person {name:"Bob"}) ,  
(sally:Person {name:"Sally"}) ,  
(anna:Person {name:"Anna"}) ,  
(jim:Person {name:"Jim"}) ,  
(mike:Person {name:"Mike"}) ,  
(billy:Person {name:"Billy"}) ,
```


```
(joe)-[:KNOWS]->(bob) ,  
(joe)-[:KNOWS]->(sally) ,  
(bob)-[:KNOWS]->(sally) ,  
(sally)-[:KNOWS]->(anna) ,  
(anna)-[:KNOWS]->(jim) ,  
(anna)-[:KNOWS]->(mike) ,  
(jim)-[:KNOWS]->(mike) ,  
(jim)-[:KNOWS]->(billy)
```



# Practical Cypher

## Social Graph - Friends of Joe's Friends

```
MATCH (person)-[:KNOWS]-(friend),  
      (friend)-[:KNOWS]-(foaf)  
WHERE person.name = "Joe"  
  AND NOT(person-[:KNOWS]-foaf)  
RETURN foaf
```




foaf

{name:"Anna"}

# Practical Cypher

## Social Graph - Common Friends

```
MATCH (person1)-[:KNOWS]-(friend),  
      (person2)-[:KNOWS]-(friend)  
WHERE person1.name = "Joe"  
      AND person2.name = "Sally"  
RETURN friend
```



friend


```
{name:"Bob"}
```

# Practical Cypher

## Social Graph - Shortest Path

```

MATCH path = shortestPath(
  (person1)-[ :KNOWS*..6 ]-(person2)
)
WHERE person1.name = "Joe"
  AND person2.name = "Billy"
RETURN path
    
```




**path**

```

{start:"13759",
nodes: [ "13759", "13757", "13756", "13755", "13753" ],
length:4,
relationships:[ "101407", "101409", "101410", "101413" ],
end:"13753"}
```

## Background

- Online jobs and career community, providing anonymized inside information to job seekers



## Business problem

- Wanted to leverage known fact that most jobs are found through personal & professional connections
- Needed to rely on an existing source of social network data. Facebook was the ideal choice.
- End users needed to get instant gratification
- Aiming to have the best job search service, in a very competitive market


## Solution & Benefits

- First-to-market with a product that let users find jobs through their network of Facebook friends
- Job recommendations served real-time from Neo4j
- Individual Facebook graphs imported real-time into Neo4j
- Glassdoor now stores > 50% of the entire Facebook social graph
- Neo4j cluster has grown seamlessly, with new instances being brought online as graph size and load have increased

# Network Management Example

## Background

- Second largest communications company in France
- Part of Vivendi Group, partnering with Vodafone



## Business problem

- Infrastructure maintenance took one full week to plan, because of the need to model network impacts
- Needed rapid, automated “what if” analysis to ensure resilience during unplanned network outages
- Identify weaknesses in the network to uncover the need for additional redundancy
- Network information spread across > 30 systems, with daily changes to network infrastructure
- Business needs sometimes changed very rapidly

## Solution & Benefits

- Flexible network inventory management system, to support modeling, aggregation & troubleshooting
- Single source of truth (Neo4j) representing the entire network
- Dynamic system loads data from 30+ systems, and allows new applications to access network data
- Modeling efforts greatly reduced because of the near 1:1 mapping between the real world and the graph
- Flexible schema highly adaptable to changing business requirements




Industry: Web/ISV, Communications  
Use case: Network Management  
Global (U.S., France)



## Background

- World's largest provider of IT infrastructure, software & services
- HP's *Unified Correlation Analyzer (UCA)* application is a key application inside HP's OSS Assurance portfolio
- Carrier-class resource & service management, problem determination, root cause & service impact analysis
- Helps communications operators manage large, complex and fast changing networks



## Business problem

- Use network topology information to identify root problems causes on the network
- Simplify alarm handling by human operators
- Automate handling of certain types of alarms Help operators respond rapidly to network issues
- Filter/group/eliminate redundant Network Management System alarms by event correlation

## Solution & Benefits

- Accelerated product development time
- Extremely fast querying of network topology
- Graph representation a perfect domain fit
- 24x7 carrier-grade reliability with Neo4j HA clustering
- Met objective in under 6 months


# Practical Cypher

## Network Management - Create



### CREATE

```
(crm {name:"CRM"}) ,  
(dbvm {name:"Database VM"}) ,  
(www {name:"Public Website"}) ,  
(wwwvm {name:"Webserver VM"}) ,  
(srv1 {name:"Server 1"}) ,  
(san {name:"SAN"}) ,  
(srv2 {name:"Server 2"}) ,  
  
(crm)-[:DEPENDS_ON]->(dbvm) ,  
(dbvm)-[:DEPENDS_ON]->(srv2) ,  
(srv2)-[:DEPENDS_ON]->(san) ,  
(www)-[:DEPENDS_ON]->(dbvm) ,  
(www)-[:DEPENDS_ON]->(wwwvm) ,  
(wwwvm)-[:DEPENDS_ON]->(srv1) ,  
(srv1)-[:DEPENDS_ON]->(san)
```



# Practical Cypher

## Network Management - Impact Analysis

```
// Server 1 Outage
MATCH (n)<-[ :DEPENDS_ON* ]-(upstream)
WHERE n.name = "Server 1"
RETURN upstream
```




upstream
{name:"Webserver VM"}
{name:"Public Website"}

# Practical Cypher

## Network Management - Dependency Analysis

```
// Public website dependencies
MATCH (n)-[ :DEPENDS_ON* ]->(downstream)
WHERE n.name = "Public Website"
RETURN downstream
```



downstream
{name:"Database VM"}
{name:"Server 2"}
{name:"SAN"}
{name:"Webserver VM"}
{name:"Server 1"}

# Practical Cypher

## Network Management - Statistics

```
// Most depended on component
MATCH (n)<-[ :DEPENDS_ON* ]-(dependent)
RETURN n,
       count(DISTINCT dependent)
           AS dependents
ORDER BY dependents DESC
LIMIT 1
```



n	dependents
{name: "SAN"}	6

# More Graphs in the Real World

## Background

- One of the world's largest logistics carriers
- Projected to outgrow capacity of old system
- New parcel routing system
  - Single source of truth for entire network
  - B2C & B2B parcel tracking
  - Real-time routing: up to 5M parcels per day



## Business problem

- 24x7 availability, year round
- Peak loads of 2500+ parcels per second
- Complex and diverse software stack
- Need predictable performance & linear scalability
- Daily changes to logistics network: route from any point, to any point


## Solution & Benefits

- Neo4j provides the ideal domain fit:
  - a logistics network is a graph
- Extreme availability & performance with Neo4j clustering
- Hugely simplified queries, vs. relational for complex routing
- Flexible data model can reflect real-world data variance much better than relational
- “Whiteboard friendly” model easy to understand

## Cisco.com

### Background

- Cisco.com serves customer and business customers with Support Services
- Needed real-time recommendations, to encourage use of online knowledge base
- Cisco had been successfully using Neo4j for its internal master data management solution.
  - Identified a strong fit for online recommendations



### Business problem

- Call center volumes needed to be lowered by improving the efficacy of online self service
- Leverage large amounts of knowledge stored in service cases, solutions, articles, forums, etc.
- Problem resolution times, as well as support costs, needed to be lowered

### Solution & Benefits

- Cases, solutions, articles, etc. continuously scraped for cross-reference links, and represented in Neo4j
- Real-time reading recommendations via Neo4j
- Neo4j Enterprise with HA cluster
- The result: customers obtain help faster, with decreased reliance on customer support

## Background

- Europe's largest communications company
- Provider of mobile & land telephone lines to consumers and businesses, as well as internet services, television, and other services

> 236,000

Employees worldwide in 2011


50

Countries

> 58 bn. €

Revenue in 2011

## Interactive Television Programming



## Business problem


- The Fanorakel application allows fans to have an interactive experience while watching sports
- Fans can vote for referee decisions and interact with other fans watching the game
- Highly connected dataset with real-time updates
- Queries need to be served real-time on rapidly changing data
- One technical challenge is to handle the very high spikes of activity during popular games

## Solution & Benefits

- Interactive, social offering gives fans a way to experience the game more closely
- Increased customer stickiness for Deutsche Telekom
- A completely new channel for reaching customers with information, promotions, and ads
- Clear competitive advantage

## Background

- Memory Lane, Inc. was founded in 1995 and based in Seattle, Washington. Subsidiary of United Online, Inc.
- Classmates.com, operates an online yearbook that connects members in the United States and Canada with friends and acquaintances from school, work, and the military.
- Evolving toward more sophisticated social networking capability



## Business problem


- Develop new Social capabilities to help monetize Yearbook-related offerings
  - Show me all the people I know in a yearbook
  - Show me yearbooks my friends appear in most often (i.e. "Top yearbooks to look at")
  - Show me sections of a yearbook that your friends appear most in (i.e. "8 of your friends are on page 12 with the football team")
  - Show me other high schools that my friends went to (i.e. friends you made in other schools)

## Solution & Benefits

- 3-Instance Neo4j Cluster with Cache Sharding + Disaster-Recovery Cluster
- Neo4j provides 18 ms response time for the top 4 queries
- Initial graph size: 100M nodes and 600M relationships
  - People, Images, Schools, Yearbooks, Yearbook Pages
- Projected to grow to 1B nodes & 6B relationships

## Background

- Teachscape, Inc. develops online learning tools for K-12 teachers, school principals, and other instructional leaders.
- Teachscape evaluated relational as an option, considering MySQL and Oracle.
- Neo4j was selected because the graph data model provides a more natural fit for managing organizational hierarchy and access to assets.



## Business problem

- Neo4j was selected to be at the heart of a new architecture.
- The user management system, centered around Neo4j, will be used to support single sign-on, user management, contract management, and end-user access to their subscription entitlements.

## Solution & Benefits

- Domain and technology fit
  - simple domain model where the relationships are relatively complex. Secondary factors included support for transactions, strong Java support, and well-implemented Lucene indexing integration
- Speed and Flexibility
  - The business depends on being able to do complex walks quickly and efficiently. This was a major factor in the decision to use Neo4j.
- Ease of Use
  - accommodate efficient access for home-grown and commercial off-the-shelf applications, as well as ad-hoc use.
  - Extreme availability & performance with Neo4j clustering
  - Hugely simplified queries, vs. relational for complex routing
  - Flexible data model can reflect real-world data variance much better than relational
  - “Whiteboard friendly” model easy to understand

## Background

- Own several job search properties, including Internships.com, TweetMyJobs, and CareerBeam
- CareerArc Group ecosystem now serves 30 million annual visitors, 40,000 companies (including more than 50% of the Fortune 100), and over 350 universities



## Business problem

- Launching new application: “Who do you know on Facebook”
- Existing SQL Server application difficult to develop against, and slowing down as user volumes increased
- Looking for a solution that could easily manage a large amount of interrelated data

## Solution & Benefits

- Neo4j's was a natural fit for the social graph underlying CareerArcGroup's professional social network
- Significantly accelerated R&D, overcoming the agility problems the team had been facing
- Strong, reliable performance on EC2, running reliably inside of existing EC2 architecture
- Initial go-live included 10 million nodes, 30 million relationships and 235 million properties

# Background

- Hong Kong based telephony infrastructure provider (aka M800 aka Pop Media)
- Exclusive China Mobile partner for international toll-free services. SMS Hub & other offerings
- 2012 Red Herring Top 100 Global Winner



## Business problem


- Launched a new mobile communication app “Maaii” allowing consumers to communicate by voice & text (Similar to Line, Viber, Rebtel, VoxOx...)
- Needed to store & relate devices, users, and contacts
- Import phone numbers from users’ address books. Rapidly serve up contacts from central database to the mobile app
- Currently around 3M users w/200M nodes in the graph

## Solution & Benefits

- Quick transactional performance for key operations:
  - friend suggestions (“friend of friend”)
  - updating contacts, blocking calls, etc.
  - etc.
- High availability telephony app uses Neo4j clustering
- Strong architecture fit: Scala w/Neo4j embedded

## Background

- London-based ‘A-list’ home listing & booking
- Web & mobile app w/visible floor plans
- Luxury benefits & home convenience in international cities (NYC, London) 24/7 phone services, suggested eateries/events/hot spots
- Travel agency partners
- 2010 had 6 apts - 2012, = 700 & growing
- Bloomberg Business “Like living in a catalog”



## Business problem

- Needed to serve up structure & floor plan details, and related spatial uses: “where’s linen?”
- Relational database couldn’t unlock the power of data & balance features with support a
- Needed complex structural data modeling with speed, accuracy and flexibility
- Serving thousands of listings & rapidly growing

## Solution & Benefits

- 6-week implementation, integrated w/framework
- optimized workflow w/user interface = easy entry
- adds operation speed/efficiency, saves time/money
- scales easily, supports growth, strong community
- Best NoSQL alternative, w/Neo4j query performance & admin interface

## Background

- Founded in 1999. Widely considered the industry leader in patient management for discharges & referrals
- Manage patient referrals for more than 4600 health care facilities
- Connects providers, payers and suppliers via secure electronic patient-transition networks, and web-based patient management platform



## Business problem

- Satisfy complex “Graph Search” queries by discharge nurses and intake coordinators, e.g.:  
*“Find a skilled nursing facility within n miles of a given location, belonging to health care group XYZ, offering speech therapy and cardiac care, and optionally Italian language services”*
- Real-time Oracle performance not satisfactory
- New functionality called for more complexity, including granular role-based access control

## Solution & Benefits

- Fast real-time performance needs now satisfied
- Queries span multiple hierarchies, including provider graph & employee permissions graph
- Graph data model provided a strong basis for adding more dimensions to the data, such as insurance networks, service areas, and ACOs (Accountable Care Organizations)
- Some multi-page SQL statements have been turned into one simple function with Neo4j

## Background

- Clinical diagnostics company specializing in genetic carrier screening for inherited diseases
- Founded in 2008 by Harvard Business School & Harvard Medical School graduates
- Two sides of the business: Clinical and R&D
- Particularly strong in the detection of rare alleles and measuring frequency in the population



## Business problem

- Clinical data split across several operational databases that are not structured for discovery
- Needed an easy query mechanism for scientists who are not data scientists. “Graph search” for bioinformatics.
- Much in Bioinformatics remains unknown: having to specifying a schema ahead of time can range from difficult to impossible.


## Solution & Benefits

- New R&D database build atop Neo4j to support information discovery by scientists
- Lightweight web front end allows simple Cypher queries to be constructed ad hoc
- Raw VCF sequence data imported into Neo4j, along with clinical data from Oracle database
- Time to answer new questions went from days of ad-hoc information gathering to hours or minutes

# Gartner's “5 Graphs”



## Consumer Web Giants Depends on Five Graphs



# Questions ?

